Ao LIU

Ph.D. in Computer Science

📱 *+1 (518) 233-4797*
📰 *Personal Website*
✉ *aoliu.cs@gmail.com*
Google Scholar

Summary

Research Interests	Machine Learning (Recommendation Systems, Learning to Rank, Algorithm Design, etc.), Differential Privacy, Computational Social Choice, and Quantum Computation
Experience	1+ years of industry experience in machine learning (personalized recommendation) Hands-on experience in recommendation systems, LLM, and theoretical computer science
Academia	Published more than 20 papers and received more than 250 citations to date. Reviewed 50+ papers for top AI conferences or journals.

Education

01/2018 – 05/2023	**Ph.D. in Computer Science,** Rensselaer Polytechnic Institute (RPI) Thesis: *Group Decision Makings from Partial Preferences* [**Link**]	Troy, NY
08/2015 – 05/2018	**M.Eng. in Material Physics,** Rensselaer Polytechnic Institute (RPI)	Troy, NY
08/2010 – 05/2014	**B.S. in Mathematics and Physics,** Tsinghua University *Minor in Computer Technology*, and *in Academic Talent Program*	Beijing, China

Experience

07/2023 – Present	**Software Engineer at Google Core ML** Design and Research More Efficient Transformers for Recommendation Systems	Mountain View, CA

05/2022 – 08/2022	**Research Intern at Google Core ML** Project: A More Accurate Position Bias Estimator for Unbiased Learning to Rank	Mountain View, CA

○ Significant accuracy improvement ($\sim$8%) *v.s.* the sate-of-art without extra computational resources.
- Proposed a more accurate model for position bias in recommendation systems.
- Designed a new machine learning framework based on the proposed position bias model and two-tower model. (*user-modeling, algorithm design, statistics*)
- Implemented a data pipeline to extract the required features from data. (*C/C++*)
- Implemented the machine learning framework, tested on public dataset, and integrated it into Google codebase. (*Python, TensorFlow, GoogleTest*)

08/2018 – 12/2018 and 05/2019 – 08/2019	**Visiting Scholar at MIT-IBM Watson AI Lab** Project: Certifiably Robust Interpretation via Rényi Differential Privacy	Cambridge, MA

○ Significant robustness improvement ($\sim$12%) plus accuracy improvement *v.s.* the sate-of-art.
- Yes! We improved both robustness and accuracy at the same time! Note that accuracy and robustness are trade-offs in machine learning.
- Theoretically connected Rényi differential privacy and interpretation robustness.
- Designed a new robust algorithm to interpret neural networks for image classifications.
- Implemented the proposed robust algorithm and tested it on VOC2007 dataset for various properties, including robustness, accuracy, and computational efficiency. (*PyTorch, TorchRay, TensorFlow*)

○ Delivered one academic paper (on top AI journal) and two patents.

Skills

Implementation	Programming languages: *Python, C/C++, MATLAB* Tools and platforms: *TensorFlow, PyTorch, TorchRay, LaTeX, GoogleTest*
Design/ Theory	*Algorithm design, Statistics, Time/Sample-complexity analysis, User-modeling, Markov chain Mount-Carlo (MCMC), Differential privacy analysis, Robustness analysis, Model identifiability.*

Review Services

Journal	Information Sciences, TMLR, ACM ToIS, Sankhya B
Conference	NeurIPS (20,21,22&23), ICML (22,23&24), ICLR (23&24), AAAI (21&22), IJCAI-22

Selected Publications

TMLR — **Smoothed Differential Privacy [PDF]**
Ao Liu, Yu-Xiang Wang, and Lirong Xia

UAI-23 — **Accelerating Voting by Quantum Computation [PDF]**
Ao Liu, Qishen Han, Lirong Xia, and Nengkun Yu

AIJ — **Certifiably Robust Interpretation via Rényi Differential Privacy [Link] [ArXiv]**
Ao Liu, Xiaoyu Chen, Sijia Liu, Lirong Xia, and Chuang Gan
Also in proceedings of **AAAI-23 Journal Track** (oral presentation)

AAAI-23 (oral) — **Differentially Private Condorcet Voting [PDF]**
Zhechen Li, _Ao Liu_, Lirong Xia, Yongzhi Cao, and Hanpin Wang

AAAI-22 — **The Semi-Random Likelihood of Doctrinal Paradoxes [PDF]**
Ao Liu, and Lirong Xia

IJCAI-22 (oral) — **Learning Mixtures of Random Utility Models with Features from Incomplete Preferences [PDF]**
Zhibing Zhao, _Ao Liu_, and Lirong Xia

JAIR — **Learning to Design Fair and Private Voting Rules [PDF]**
Farhad Mohsin, _Ao Liu_, Pin-Yu Chen, Francesca Rossi, and Lirong Xia

UAI-20 (oral) — **How Private Are Commonly-Used Voting Rules? [PDF]**
Farhad Mohsin, _Ao Liu_, Pin-Yu Chen, Francesca Rossi, and Lirong Xia

ETRA-20 Adjunct — **Let It Snow: Adding Pixel Noise to Protect the Users Identity [Link]**
Brendan John, _Ao Liu_, Lirong Xia, Sanjeev Koppal, and Eakta Jain

AAAI-19 (oral) — **Near-Neighbor Methods in Random Preference Completion [PDF]**
Ao Liu, Qiong Wu, Zhenming Liu, and Lirong Xia

AAAI-19 (oral) — **Learning Plackett-Luce Mixture from Partial Preferences [PDF]**
Ao Liu, Zhibing Zhao, Chao Liao, Pinyan Lu, and Lirong Xia

ETRA-19 (oral) — **Differential Privacy for Eye-Tracking Data [PDF]**
Ao Liu, L. Xia, A. Duchowski, R. Bailey, K. Holmqvist, and E. Jain

US Patent — **Certifiably Robust Interpretation [PDF]**
Ao Liu, Sijia Liu, Bo Wu, Lirong Xia, Qi Cheng Li, and Chuang Gan

US Patent — **Interpretation Maps with Guaranteed Robustness [PDF]**
Ao Liu, Sijia Liu, Abhishek Bhandwaldar, Chuang Gan, Lirong Xia, and Qi Cheng Li

Awards and Teaching

09/2019 – 05/2022 — **RPI-IBM AI Horizon Scholarship**
3-year scholarship supported by Rensselaer-IBM Artificial Intelligence Research Collaboration

09/2016 – 05/2017 — **RPI Presidential Graduate Research Fellowship**
A One-Year Fellowship for Outstanding Graduate Students [**Certificate**]

Spring 2023 — **Teaching Assistant** of **CSCI 4150: Introduction to AI**
Instructor: Lirong Xia

04/2021 — **Guest Lecture** at **CSCI 4967/6967: Economics and Computation**
Topic: The Semi-Random Likelihood of Doctrinal Paradoxes

Spring 2023 — **Teaching Assistant** of **MATH 1020: Calculus II**
Instructor: David A. Schmidt